



P-PAS Makes Nemotron-3.5-Lightning-30B Faster

P-PAS dynamically adapts the vLLM scheduling budget as serving pressure changes.



P-PAS achieves the lowest average latency at every tested burst rate



One fixed scheduling budget does not fit every serving regime

Low pressure favors large budgets, while high pressure favors much smaller ones. P-PAS adapts between them dynamically.

vLLM 0.27.1 • P-PAS • NVIDIA Nemotron-3.5-Lightning-30B-A3B-NVFP4 • RTX 5090